

```
In [1]: ## Rebecca Lewis  
## DSC 650  
## Assignment 1.1
```

```
In [1]: '''Trains a simple deep NN on the MNIST dataset.

Gets to 98.40% test accuracy after 20 epochs
(there is *a lot* of margin for parameter tuning).
2 seconds per epoch on a K520 GPU.
'''

from tensorflow import keras
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import RMSprop

batch_size = 128
num_classes = 10
epochs = 20

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>)

11493376/11490434 [=====] - 1s 0us/step  
 60000 train samples  
 10000 test samples  
 Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401920
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 512)	262656
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130

Total params: 669,706  
 Trainable params: 669,706  
 Non-trainable params: 0

Epoch 1/20

469/469 [=====] - 5s 10ms/step - loss: 0.2471 - accuracy: 0.9237 - val\_loss: 0.1100 - val\_accuracy: 0.9650

Epoch 2/20

469/469 [=====] - 5s 10ms/step - loss: 0.1022 - accuracy: 0.9689 - val\_loss: 0.0836 - val\_accuracy: 0.9728

Epoch 3/20

469/469 [=====] - 4s 10ms/step - loss: 0.0742 - accuracy: 0.9774 - val\_loss: 0.0774 - val\_accuracy: 0.9773

Epoch 4/20

469/469 [=====] - 4s 9ms/step - loss: 0.0609 - accuracy: 0.9813 - val\_loss: 0.0752 - val\_accuracy: 0.9786

Epoch 5/20

469/469 [=====] - 4s 9ms/step - loss: 0.0512 - accuracy: 0.9838 - val\_loss: 0.0715 - val\_accuracy: 0.9803

Epoch 6/20

469/469 [=====] - 4s 9ms/step - loss: 0.0423 - accuracy: 0.9876 - val\_loss: 0.0830 - val\_accuracy: 0.9794

Epoch 7/20

469/469 [=====] - 4s 9ms/step - loss: 0.0395 - accuracy: 0.9883 - val\_loss: 0.0844 - val\_accuracy: 0.9800

Epoch 8/20

469/469 [=====] - 4s 9ms/step - loss: 0.0325 - accuracy: 0.9902 - val\_loss: 0.0946 - val\_accuracy: 0.9812

Epoch 9/20

469/469 [=====] - 4s 9ms/step - loss: 0.0299 - accuracy: 0.9909 - val\_loss: 0.0821 - val\_accuracy: 0.9819

Epoch 10/20

469/469 [=====] - 4s 9ms/step - loss: 0.0274 - accuracy: 0.9914 - val\_loss: 0.0821 - val\_accuracy: 0.9832

Epoch 11/20

469/469 [=====] - 4s 9ms/step - loss: 0.0250 - accuracy: 0.9926 - val\_loss: 0.0939 - val\_accuracy: 0.9823

```
Epoch 12/20
469/469 [=====] - 4s 9ms/step - loss: 0.0264 - accur
acy: 0.9930 - val_loss: 0.0878 - val_accuracy: 0.9834
Epoch 13/20
469/469 [=====] - 4s 9ms/step - loss: 0.0224 - accur
acy: 0.9935 - val_loss: 0.0903 - val_accuracy: 0.9849
Epoch 14/20
469/469 [=====] - 4s 9ms/step - loss: 0.0227 - accur
acy: 0.9933 - val_loss: 0.0980 - val_accuracy: 0.9832
Epoch 15/20
469/469 [=====] - 4s 9ms/step - loss: 0.0190 - accur
acy: 0.9944 - val_loss: 0.1142 - val_accuracy: 0.9827
Epoch 16/20
469/469 [=====] - 4s 9ms/step - loss: 0.0210 - accur
acy: 0.9944 - val_loss: 0.0944 - val_accuracy: 0.9837
Epoch 17/20
469/469 [=====] - 4s 9ms/step - loss: 0.0166 - accur
acy: 0.9949 - val_loss: 0.1246 - val_accuracy: 0.9834
Epoch 18/20
469/469 [=====] - 4s 9ms/step - loss: 0.0170 - accur
acy: 0.9948 - val_loss: 0.1206 - val_accuracy: 0.9827
Epoch 19/20
469/469 [=====] - 4s 9ms/step - loss: 0.0152 - accur
acy: 0.9956 - val_loss: 0.1246 - val_accuracy: 0.9830
Epoch 20/20
469/469 [=====] - 4s 9ms/step - loss: 0.0179 - accur
acy: 0.9953 - val_loss: 0.1278 - val_accuracy: 0.9841
Test loss: 0.12781260907649994
Test accuracy: 0.9840999841690063
```

```

In [9]: #
# Licensed to the Apache Software Foundation (ASF) under one or more
# contributor license agreements. See the NOTICE file distributed with
# this work for additional information regarding copyright ownership.
# The ASF licenses this file to You under the Apache License, Version 2.0
# (the "License"); you may not use this file except in compliance with
# the License. You may obtain a copy of the License at
#
# http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
#

import sys
from random import random
from operator import add

from pyspark.sql import SparkSession

if __name__ == "__main__":
    """
        Usage: pi [partitions]
    """
    spark = SparkSession\
        .builder\
        .appName("PythonPi")\
        .getOrCreate()

    #receiving a base 10 error with original code, added a base of 16 but was get
    #for a negative value. Applied the absolute value and was able to run succes
    partitions = abs(int(sys.argv[1], base=16)) if len(sys.argv) > 1 else 2
    n = 100000 * partitions

    def f(_):
        x = random() * 2 - 1
        y = random() * 2 - 1
        return 1 if x ** 2 + y ** 2 <= 1 else 0

    count = spark.sparkContext.parallelize(range(1, n + 1), partitions).map(f).re
    print("Pi is roughly %f" % (4.0 * count / n))

    spark.stop()

```

Pi is roughly 3.138987

In [7]: `print (sys.argv)`

```
['/opt/conda/lib/python3.8/site-packages/ipykernel_launcher.py', '-f', '/home/jovyan/.local/share/jupyter/runtime/kernel-5dabe6bd-6159-4d02-af0e-b7446b39bf31.json']
```

In [ ]: